

ScholarSphere - AI-Powered Research Paper Collaboration Hub

Updated for Postgres + pgvector, Next.js (App Router) frontend, and a separate Node.js + Express backend with Prisma. This document outlines the architecture, features, and implementation details for ScholarSphere, a SaaS platform designed to help researchers manage academic papers, collaborate with teams, and leverage AI for enhanced workflows.

1) Product Overview

- Name: ScholarSphere
- Type: SaaS platform for research paper organization, collaboration, and AI assistance
- Audience: Researchers, students, professors, academic teams, institutions
- Goal: Help users upload, manage, annotate, and collaborate on academic papers; power workflows with AI (summaries, semantic search, citation tooling, literature review assistance); provide team workspaces and paid plans.

2) System Architecture (Separated Frontend and Backend)

High level

- Frontend Web App (Next.js, Vercel):
 - Next.js App Router (React, TypeScript, Tailwind, ShadCN).
 - Auth handled by Auth.js (NextAuth) with JWT sessions.
 - Calls the backend API with Bearer JWT; uses RTK Query for data fetching.
- Backend API (Node.js, Express, Prisma, Postgres):
 - A dedicated service (Railway/Render/Fly.io). Single source of truth for business logic.
 - Validates Auth.js JWT via shared NEXTAUTH_SECRET (HS256) or JWKS if using asymmetric signing.
 - Exposes REST endpoints for papers, annotations, search, collections, workspaces, billing.
 - Handles payment webhooks (Stripe/SSLCommerz) and writes Payment/Subscription records.
- Worker/Jobs Service (optional but recommended):
 - Handles heavy tasks: file parsing, OCR, chunking, embeddings, AI pipelines.
 - Queue with Redis + BullMQ. Workers consume jobs and persist results via Prisma.
- Database & Storage:
 - Postgres 15+ with pgvector extension; accessed via Prisma.
 - Object storage (S3-compatible) for PDFs and assets; signed URLs from backend.
 - CDN in front of storage for fast delivery.
- Observability & Operations:
 - PostHog/Amplitude for product analytics; OpenTelemetry/logging for API observability.
 - Alerting on queue backlogs, failed webhooks, and AI provider errors.

Data flow highlights

- Upload: Frontend uploads file -> Backend pre-signs S3 URL -> Client uploads -> Backend enqueues "ingest" job -> Worker parses/OCRs -> chunks + embeddings -> records in DB -> UI shows progress.
- Auth: Frontend signs in via Auth.js -> stores JWT -> sends to Backend in Authorization header -> Backend enforces RBAC.
- Billing: Frontend starts checkout -> redirects to provider -> provider posts webhook to Backend -> Backend activates Subscription -> notifies Frontend via polling or SSE/Webhook to client.

3) Core Features

A. Paper Management

- Upload PDFs/DOCX/LaTeX, import via DOI/arXiv/OpenAlex/Semantic Scholar.
- OCR for scanned docs (Tesseract); parse text; extract title/authors/abstract via regex + NLP.
- Smart tagging (topics/methods) via OpenAI/HF; short AI summaries.
- Bulk import with progress tracking.

B. Citation & References

- Citation formatting (APA/MLA/IEEE).
- Citation graph visualization; missing-citation suggestions (LLM + similarity).
- Export citations; per-collection bibliography.

C. Semantic Search & Discovery

- Chunked text (500–1,000 tokens, 50–100 overlap), embeddings per chunk (1536 or 3072 dims).
- pgvector with ivfflat index and cosine distance; filters by workspace, year, tags.
- Similar papers recommender; trends via keyword extraction.

D. Collaboration

- Annotations: highlights, comments, notes; threading; version history and revert.
- Collections per workspace with membership; activity logs; sharing controls.
- Real-time comments/annotation updates via WebSocket/SSE.

E. AI Writing Assistance

- Abstract generator, literature review outliner.
- Self-plagiarism check via cosine similarity across user's papers.

F. Payments & Plans

- Stripe (global), SSLCommerz (BD). Freemium + Pro + Institutional.
- Webhooks update Payments/Subscriptions; customer portal links for Stripe.

4) Detailed Tech Stack

- Frontend: Next.js (App Router), React, TypeScript, Tailwind, ShadCN, Framer Motion; Forms with RHF + Zod; State with RTK Query + light UI slice.
- Backend API: Node.js, Express.js, Prisma ORM; Zod for input validation; Helmet, CORS; RBAC middleware by role and workspace membership.
- Workers: Node.js with BullMQ (Redis) for ingest/OCR/embeddings/AI tasks.
- Database: Postgres + Prisma; pgvector for embeddings (stored as Unsupported("vector") in Prisma, queried via raw SQL).
- AI: OpenAI, HuggingFace, LangChain, Tesseract.
- Storage: S3-compatible (R2/S3/MinIO) with signed URLs.
- Payments: Stripe SDK; SSLCommerz (sslcommerz-lts).
- DevOps: Vercel (frontend); Railway/Render/Fly.io (backend + workers); Redis for queues/cache.

5) Authentication & Authorization

- Auth.js (NextAuth) on the frontend app; JWT session strategy.
- Backend validates JWT (bearer) with shared secret; extracts userId and role.
- Workspace-based access: WorkspaceMember role gates access to papers/collections inside a workspace.
- Collection-level sharing for granular collaboration.
- Admin users can manage users, payments, and system settings.

6) API Surface (Backend Express)

- Auth
 - POST /auth/session/validate (optional ping) – verifies JWT valid/claims
- Papers
 - POST /papers/upload-url – get pre-signed upload URL
 - POST /papers/import – import by DOI/API providers
 - GET /papers – list/filter; supports semantic=true for vector search
 - GET /papers/:id – detail with metadata
 - DELETE /papers/:id – soft delete
- Annotations
 - GET /papers/:id/annotations
 - POST /papers/:id/annotations
 - GET /papers/:id/annotations/versions
 - POST /papers/:id/annotations/versions/revert
- Collections
 - GET /collections
 - POST /collections
 - GET /collections/:id
 - POST /collections/:id/papers
 - POST /collections/:id/invite
- Search & AI
 - POST /search/semantic
 - GET /papers/:id/similar
 - POST /papers/:id/ai/summarize
 - POST /papers/:id/ai/citation-suggestions
- Graph
 - GET /graph/paper/:id
- Workspaces
 - POST /workspaces
 - POST /workspaces/:id/invite
 - PUT /workspaces/:id/members/:userId
 - GET /workspaces/:id/activities
- Billing
 - GET /subscriptions
 - POST /subscriptions/checkout
 - GET /payments
 - POST /payments/ssl/init
 - POST /webhooks/stripe
 - POST /webhooks/sslcommerz
- Admin
 - GET /admin/users
 - PUT /admin/users/:id
 - GET /admin/metrics

Security and cross-cutting

- CORS: allow frontend domain; block others by default.
- CSRF: not required for pure bearer APIs; ensure cookie usage is httpOnly/secure if used.
- Rate limiting: per-IP and per-user; stricter on AI endpoints.
- Input validation: Zod schemas on every route; sanitize filters for search.
- Audit logs and notifications for sensitive actions.

7) UI/UX – Page-by-Page (Frontend Next.js)

Global Chrome

- Sidebar: Dashboard, Papers, Collections, Upload, Graph, Teams, Billing, Admin
- Header: Global search, notifications, quick actions, profile menu
- Mobile FAB: Upload / New annotation

1. Landing (Marketing)

- Hero CTA, features, pricing, testimonials; analytics events.
2. Auth
 - OAuth (Google/ORCID/GitHub) or email magic link; profile completion (name, institution, role).
 3. Onboarding
 - Upload/import first paper; create workspace and collection; invite teammates; guided tips.
 4. Dashboard
 - Recent uploads, recommended papers, active collections, notifications; quick actions.
 5. Papers Library
 - Search bar with keyword/semantic toggle; filters; list/grid; bulk actions (add to collection, export citations).
 6. Paper Detail / Reader
 - PDF viewer + extracted-text toggle; annotation panel (threaded); AI actions (summarize, cite, paraphrase); versioning and revert.
 7. Collections
 - List + detail (papers grid, activity, access control); drag-and-drop; invites; export.
 8. Semantic Search / Discovery
 - Ranked results with similarity score; snippet previews; similar-papers panel; trends.
 9. Citation Graph
 - Interactive graph (D3/Cytoscape); cluster/filter; node panel with open/cite actions.
 10. Upload / Import
 - Drag-drop; OCR toggle; import by DOI/URL or providers; background job status.
 11. Annotation History
 - Version timeline; side-by-side diff; restore.
 12. Team / Workspace
 - Workspace list and details; members and roles; invites; activity log.
 13. Billing
 - Plan status; upgrade/checkout; payment history; invoices; customer portal link.
 14. Admin
 - User management; payments analytics (MRR, churn); system logs; API keys & quotas.
 15. Settings / Profile
 - Profile, security, integrations (ORCID, Scholar), BYO OpenAI key, data export/delete.
 16. Help / Docs
 - FAQs, tutorials, support tickets; in-app tooltips.

8) Non-Functional Requirements

- Performance: paginate lists; stream large PDFs; index pgvector with ivfflat (cosine).
- Reliability: idempotent payment webhooks; store raw webhook payloads; retries with backoff.
- Background jobs: all heavy tasks (OCR/embeddings) offloaded; job results written back atomically.
- Security: signed URLs for file access; strict RBAC; input validation; secrets management (runtime env).
- Internationalization: i18n-ready; date-fns for formatting.
- Accessibility: keyboard-friendly annotations; ARIA for viewer controls.

9) Monetization

- Free: up to 100 papers, basic AI features, no shared collections.
- Pro (\$10/mo): unlimited uploads, all AI tools, collaboration, priority queue.
- Institutional: workspace/institution-wide seats and SSO (future).

10) Prisma Schema (without ERD generator config)

The canonical source for the ERD. This schema uses UUID string IDs and includes createdAt, updatedAt, isDeleted on every model.

```
// Prisma schema for RefAIra (Postgres + pgvector)
// - All ids are String UUIDs
// - Every model includes createdAt, updatedAt, isDeleted

generator client {
  provider = "prisma-client-js"
```

```

    provider      = prisma-client-js
    previewFeatures = ["postgresqlExtensions"]
  }

datasource db {
  provider = "postgresql"
  url      = env("DATABASE_URL")
}

enum Role {
  RESEARCHER
  PRO_RESEARCHER
  TEAM_LEAD
  ADMIN
}

enum AnnotationType {
  HIGHLIGHT
  COMMENT
  NOTE
}

enum PaymentProvider {
  STRIPE
  SSLCOMMERZ
}

enum PaymentStatus {
  PENDING
  SUCCEEDED
  FAILED
  REFUNDED
}

enum SubscriptionStatus {
  ACTIVE
  EXPIRED
  CANCELED
  PAST_DUE
}

enum PlanTier {
  FREE
  PRO
  INSTITUTIONAL
}

model User {
  id      String @id @default(uuid())
  email   String @unique
  name    String?
  image   String?
  role    Role   @default(RESEARCHER)

  // Auth.js relations
  accounts Account[]
  sessions Session[]

  // App relations
  memberships      WorkspaceMember[]
  workspacesOwned  Workspace[]      @relation("WorkspaceOwner")
  uploadedPapers   Paper[]          @relation("PaperUploader")
  annotations       Annotation[]
  collections       Collection[]     @relation("CollectionOwner")
  subscriptions     Subscription[]
  payments           Payment[]

```

```

searchHistory SearchHistory[]
notifications Notification[]
usageEvents UsageEvent[]
activities ActivityLog[]

createdAt DateTime @default(now())
updatedAt DateTime @updatedAt
isDeleted Boolean @default(false)
AnnotationVersion AnnotationVersion[]
CollectionPaper CollectionPaper[]
CollectionMember CollectionMember[]
}

model Workspace {
  id String @id @default(uuid())
  name String
  ownerId String
  owner User @relation("WorkspaceOwner", fields: [ownerId], references: [id])
  members WorkspaceMember[]
  collections Collection[]
  papers Paper[]
  activities ActivityLog[]

  createdAt DateTime @default(now())
  updatedAt DateTime @updatedAt
  isDeleted Boolean @default(false)
  Subscription Subscription[]
  UsageEvent UsageEvent[]

  @@index([ownerId])
}

model WorkspaceMember {
  id String @id @default(uuid())
  workspaceId String
  userId String
  role Role @default(RESEARCHER)
  joinedAt DateTime @default(now())

  workspace Workspace @relation(fields: [workspaceId], references: [id])
  user User @relation(fields: [userId], references: [id])

  createdAt DateTime @default(now())
  updatedAt DateTime @updatedAt
  isDeleted Boolean @default(false)

  @@unique([workspaceId, userId])
  @@index([userId])
}

model Paper {
  id String @id @default(uuid())
  workspaceId String
  uploaderId String
  title String
  abstract String?
  metadata Json @db.JsonB
  source String? // 'upload' | 'arxiv' | 'openalex' | 'doi' | 'semantic-scholar'
  doi String? @unique

  uploader User @relation("PaperUploader", fields: [uploaderId], references: [id])
  workspace Workspace @relation(fields: [workspaceId], references: [id])
  file PaperFile?
  chunks PaperChunk[]
  citationsFrom Citation[] @relation("CitationsFrom")

```

```

citationsTo      Citation[]      @relation("CitationsTo")
annotations      Annotation[]
collectionJoins  CollectionPaper[]
aiSummaries      AISummary[]

createdAt  DateTime  @default(now())
updatedAt  DateTime  @updatedAt
isDeleted  Boolean   @default(false)
UsageEvent UsageEvent[]

@@index([workspaceId, createdAt])
@@index([uploaderId])
}

model PaperFile {
  id          String  @id @default(uuid())
  paperId     String  @unique
  storageProvider String // s3 | gcs | local
  objectKey   String
  contentType  String?
  sizeBytes   Int?
  pageCount   Int?
  checksum    String?

  paper Paper @relation(fields: [paperId], references: [id])

  createdAt DateTime @default(now())
  updatedAt DateTime @updatedAt
  isDeleted Boolean  @default(false)
}

model PaperChunk {
  id          String  @id @default(uuid())
  paperId     String
  idx         Int
  page        Int?
  content     String
  embedding   Unsupported("vector")? // pgvector column (e.g., vector(1536)); ERD does not depend on actual DB
  tokenCount  Int?

  paper Paper @relation(fields: [paperId], references: [id])

  createdAt DateTime @default(now())
  updatedAt DateTime @updatedAt
  isDeleted Boolean  @default(false)

  @@unique([paperId, idx])
  @@index([paperId])
}

model Citation {
  id          String  @id @default(uuid())
  sourcePaperId String
  targetPaperId String
  context     String?
  location    String?

  sourcePaper Paper @relation("CitationsFrom", fields: [sourcePaperId], references: [id])
  targetPaper  Paper @relation("CitationsTo", fields: [targetPaperId], references: [id])

  createdAt DateTime @default(now())
  updatedAt DateTime @updatedAt
  isDeleted Boolean  @default(false)

  @@unique([sourcePaperId, targetPaperId, context])
  @@index([targetPaperId])
}

```

```

    @id @default(uuid())
}

model Annotation {
    id          String      @id @default(uuid())
    paperId     String
    userId      String
    type        AnnotationType @default(HIGHLIGHT)
    anchor      Json        @db.JsonB // position anchors
    text        String
    version     Int          @default(1)
    parentId    String?

    parent      Annotation? @relation("AnnotationThread", fields: [parentId], references: [id])
    children    Annotation[] @relation("AnnotationThread")

    paper       Paper        @relation(fields: [paperId], references: [id])
    user        User         @relation(fields: [userId], references: [id])
    versions    AnnotationVersion[]

    createdAt   DateTime @default(now())
    updatedAt   DateTime @updatedAt
    isDeleted   Boolean   @default(false)

    @@index([paperId])
    @@index([userId])
}

model AnnotationVersion {
    id          String      @id @default(uuid())
    annotationId String
    version     Int
    text        String
    changedById String
    timestamp   DateTime @default(now())

    annotation Annotation @relation(fields: [annotationId], references: [id])
    changedBy  User       @relation(fields: [changedById], references: [id])

    createdAt   DateTime @default(now())
    updatedAt   DateTime @updatedAt
    isDeleted   Boolean   @default(false)

    @@unique([annotationId, version])
}

model Collection {
    id          String      @id @default(uuid())
    workspaceId String
    ownerId     String
    name        String
    description  String?
    isPublic    Boolean @default(false)

    owner       User        @relation("CollectionOwner", fields: [ownerId], references: [id])
    workspace   Workspace @relation(fields: [workspaceId], references: [id])

    papers      CollectionPaper[]
    members     CollectionMember[]

    createdAt   DateTime @default(now())
    updatedAt   DateTime @updatedAt
    isDeleted   Boolean   @default(false)

    @@index([workspaceId])
    @@index([ownerId])
}

```

```

}

model CollectionPaper {
  id          String  @id @default(uuid())
  collectionId String
  paperId     String
  addedById   String
  addedAt     DateTime @default(now())

  collection Collection @relation(fields: [collectionId], references: [id])
  paper      Paper      @relation(fields: [paperId], references: [id])
  addedBy    User        @relation(fields: [addedById], references: [id])

  createdAt DateTime @default(now())
  updatedAt DateTime @updatedAt
  isDeleted Boolean  @default(false)

  @@unique([collectionId, paperId])
  @@index([paperId])
}

```

```

model CollectionMember {
  id          String  @id @default(uuid())
  collectionId String
  userId       String
  role         Role    @default(RESEARCHER)
  invitedAt    DateTime @default(now())

  collection Collection @relation(fields: [collectionId], references: [id])
  user      User        @relation(fields: [userId], references: [id])

  createdAt DateTime @default(now())
  updatedAt DateTime @updatedAt
  isDeleted Boolean  @default(false)

  @@unique([collectionId, userId])
}

```

```

model SearchHistory {
  id      String @id @default(uuid())
  userId  String
  query   String
  filters Json?  @db.JsonB
  results Json?  @db.JsonB

  user User @relation(fields: [userId], references: [id])

  createdAt DateTime @default(now())
  updatedAt DateTime @updatedAt
  isDeleted Boolean  @default(false)

  @@index([userId, createdAt])
}

```

```

model AISummary {
  id          String @id @default(uuid())
  paperId     String
  model       String
  summary     String
  promptHash  String

  paper Paper @relation(fields: [paperId], references: [id])

  createdAt DateTime @default(now())
  updatedAt DateTime @updatedAt
  isDeleted Boolean  @default(false)
}

```



```

isDeleted Boolean @default(false)

@@unique([paperId, model, promptHash])
}

model Plan {
  id          String @id @default(uuid())
  code        String @unique
  name        String
  priceCents  Int
  currency    String
  interval    String // month, year, etc.
  stripePriceId String?
  features    Json? @db.JsonB
  active      Boolean @default(true)

  subscriptions Subscription[]

  createdAt DateTime @default(now())
  updatedAt DateTime @updatedAt
  isDeleted Boolean @default(false)
}

model Subscription {
  id          String @id @default(uuid())
  userId      String
  workspaceId String?
  planId      String
  status      SubscriptionStatus
  provider     PaymentProvider
  providerCustomerId String?
  providerSubscriptionId String?
  cancelAtPeriodEnd Boolean @default(false)
  startedAt   DateTime @default(now())
  expiresAt   DateTime

  user      User @relation(fields: [userId], references: [id])
  workspace Workspace? @relation(fields: [workspaceId], references: [id])
  plan      Plan @relation(fields: [planId], references: [id])
  payments  Payment[]

  createdAt DateTime @default(now())
  updatedAt DateTime @updatedAt
  isDeleted Boolean @default(false)

  @@index([userId])
  @@index([workspaceId])
  @@index([planId])
}

model Payment {
  id          String @id @default(uuid())
  userId      String
  subscriptionId String?
  provider     PaymentProvider
  amountCents  Int
  currency    String
  transactionId String @unique
  status       PaymentStatus
  raw          Json? @db.JsonB

  user      User @relation(fields: [userId], references: [id])
  subscription Subscription? @relation(fields: [subscriptionId], references: [id])

  createdAt DateTime @default(now())
  updatedAt DateTime @updatedAt
}

```

```

isDeleted Boolean @default(false)

@@index([userId, createdAt])
}

model WebhookEvent {
  id String @id @default(uuid())
  provider PaymentProvider
  eventId String
  type String
  payload Json @db.JsonB
  receivedAt DateTime @default(now())
  processedAt DateTime?
  status String
  error String?

  createdAt DateTime @default(now())
  updatedAt DateTime @updatedAt
  isDeleted Boolean @default(false)

  @@unique([provider, eventId])
}

model Notification {
  id String @id @default(uuid())
  userId String
  type String
  payload Json @db.JsonB
  readAt DateTime?

  user User @relation(fields: [userId], references: [id])

  createdAt DateTime @default(now())
  updatedAt DateTime @updatedAt
  isDeleted Boolean @default(false)

  @@index([userId, createdAt])
}

model UsageEvent {
  id String @id @default(uuid())
  userId String
  workspaceId String?
  kind String // "upload", "ai_summarize", "semantic_search", etc.
  units Int @default(1)
  paperId String?

  user User @relation(fields: [userId], references: [id])
  workspace Workspace? @relation(fields: [workspaceId], references: [id])
  paper Paper? @relation(fields: [paperId], references: [id])

  createdAt DateTime @default(now())
  updatedAt DateTime @updatedAt
  isDeleted Boolean @default(false)

  @@index([userId, createdAt])
}

model ActivityLog {
  id String @id @default(uuid())
  userId String?
  workspaceId String?
  entity String
  entityId String
  action String

```

```

details      Json?      @db.JsonB

user         User?       @relation(fields: [userId], references: [id])
workspace    Workspace? @relation(fields: [workspaceId], references: [id])

createdAt    DateTime    @default(now())
updatedAt    DateTime    @updatedAt
isDeleted    Boolean     @default(false)

@@index([workspaceId, createdAt])
}

// ----- Auth.js (NextAuth) models -----

model Account {
  id          String      @id @default(uuid())
  userId       String
  type         String
  provider     String
  providerAccountId String
  refresh_token String?   @db.Text
  access_token  String?   @db.Text
  expires_at    Int?
  token_type    String?
  scope         String?
  id_token      String?   @db.Text
  session_state String?

  user User @relation(fields: [userId], references: [id], onDelete: Cascade)

  createdAt DateTime @default(now())
  updatedAt DateTime @updatedAt
  isDeleted Boolean  @default(false)

  @@unique([provider, providerAccountId])
  @@index([userId])
}

model Session {
  id          String      @id @default(uuid())
  sessionToken String      @unique
  userId       String
  expires      DateTime

  user User @relation(fields: [userId], references: [id], onDelete: Cascade)

  createdAt DateTime @default(now())
  updatedAt DateTime @updatedAt
  isDeleted Boolean  @default(false)
}

model VerificationToken {
  id          String      @id @default(uuid())
  identifier String
  token       String      @unique
  expires     DateTime

  createdAt DateTime @default(now())
  updatedAt DateTime @updatedAt
  isDeleted Boolean  @default(false)

  @@unique([identifier, token])
}

```

11) Deployment

- Frontend: Vercel project with environment-bound NEXTAUTH_SECRET, NEXT_PUBLIC_API_BASE_URL.
- Backend: Railway/Render/Fly.io container; set DATABASE_URL, JWT/Auth secrets, Stripe/SSL keys, S3 creds.
- Workers: Same platform as backend, separate process using the same codebase (monorepo) or separate repo.
- Migrations: Prisma migrate on backend deploy; maintain seed scripts for local/dev environments.

12) Next Steps

- Implement Auth.js JWT strategy; add backend middleware to validate tokens and attach user context.
- Build upload flow with S3 pre-signed URLs and background ingest job.
- Implement vector search queries using pgvector + raw SQL; add ivfflat index.
- Wire Stripe/SSLCommerz checkout and webhooks; create Subscription/Payment write paths.
- Harden RBAC at workspace and collection boundaries; add activity logging for sensitive flows.

If you want, I can generate a base Express project structure (controllers, routers, middleware, Prisma client, BullMQ setup) and a Next.js starter wired to this API.